

AI VIET NAM

@aivietnam.edu.vn

Linear Regression *(Review and Assignments)*

Vinh Dinh Nguyen
PhD in Computer Science

Outline



➤ **Linear Regression Review**

➤ **Exercise 1**

➤ **Exercise 2**

➤ **Exercise 3**

➤ **Exercise 4**

➤ **Exercise 5**

➤ **Other Discussions**

MACHINE LEARNING

SUPERVISED LEARNING

Predictive Analysis
and Forecasting




SEMI- SUPERVISED LEARNING

Hybrid modeling
with labeled and
unlabeled data



UNSUPERVISED LEARNING


Raw inferences and
Pattern finding



REINFORCEMENT LEARNING

Learn from Mistakes
and old data





Linear Regression: Quick Review

Area (feet ²)	Price (\$)
2140	460k
1416	232k
1534	315k
852	178k
500	135k

Training Set



Learning Algorithm



Size of
house



h




Estimated
price


How do we represent h ?

Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$

Linear Regression with One Variable
Univariate linear regression

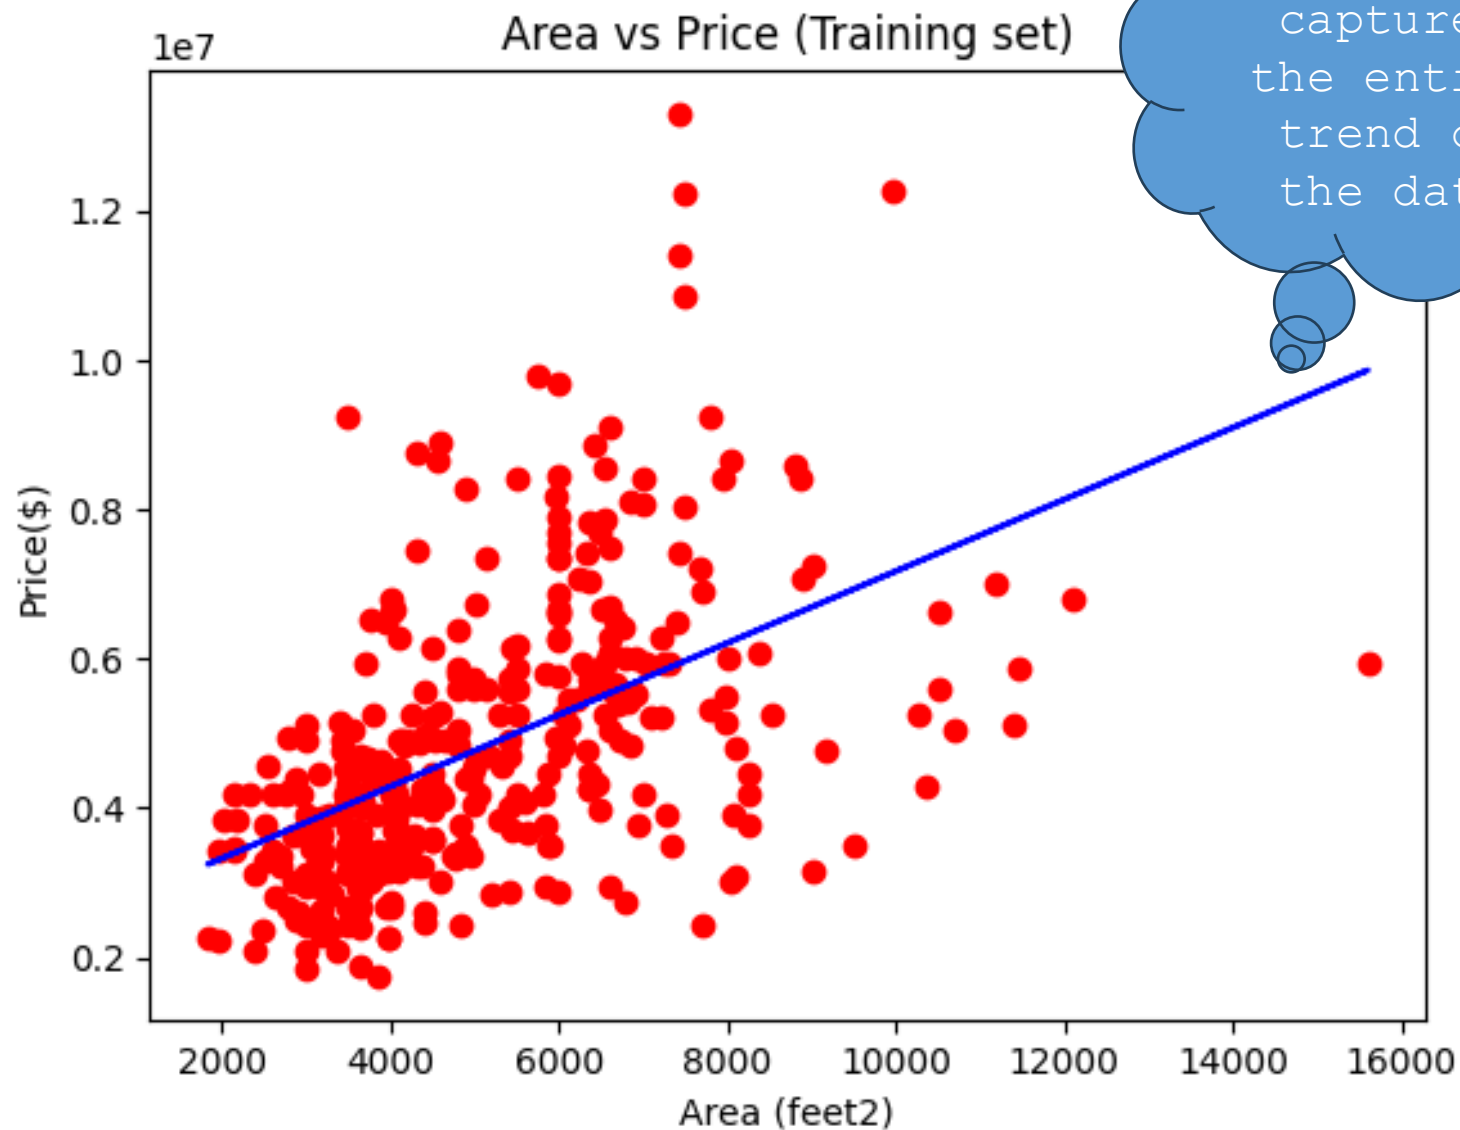


Linear Regression: Quick Review



Linear Regression

Can we find
a line that
captures
the entire
trend of
the data



Linear Regression: Quick Review

Training Set	Size in feet ² (x)	Price (\$) in 1000's (y)
	2104	460
	1416	232
	1534	315
	852	178


Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$

θ_i 's: Parameters


How to choose θ_i 's ?

Linear Regression: Quick Review


$$h_{\theta}(x) = \theta_0 + \theta_1 x$$



$$\begin{aligned}\theta_0 &= 1.5 \\ \theta_1 &= 0\end{aligned}$$




$$\begin{aligned}\theta_0 &= 0 \\ \theta_1 &= 0.5\end{aligned}$$



$$\begin{aligned}\theta_0 &= 1 \\ \theta_1 &= 0.5\end{aligned}$$

Linear Regression: Quick Review



$$J(\theta) = \frac{1}{2m} \sum (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Idea: Choose θ_0, θ_1 so that $h_{\theta}(x)$ is close to y for our training examples (x, y)

Linear Regression: Quick Review

Hypothesis:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Parameters:

$$\theta_0, \theta_1$$

Cost Function:

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Goal: minimize $J(\theta_0, \theta_1)$
 θ_0, θ_1

Simplified

$$h_{\theta}(x) = \theta_1 x$$


$$\theta_1$$

$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

minimize $J(\theta_1)$
 θ_1

$$h_{\theta}(x)$$


(for fixed θ_1 , this is a function of x)



$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$J(\theta_1)$$


(function of the parameter θ_1)



$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$


$h_{\theta}(x)$

(for fixed θ_1 , this is a function of x)



$J(\theta_1)$


(function of the parameter θ_1)



$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$


$h_{\theta}(x)$

(for fixed θ_1 , this is a function of x)



$J(\theta_1)$

(function of the parameter θ_1)



Linear Regression: Quick Review

Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$


Parameters: θ_0, θ_1

Cost Function: $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

Goal: $\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1)$

$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$h_{\theta}(x) = 50 + 0.06x$$

$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)


Linear Regression: Quick Review


Have some function $J(\theta_0, \theta_1)$


Want $\min_{\theta_0, \theta_1} J(\theta_0, \theta_1)$

Outline:

- Start with some θ_0, θ_1
- Keep changing θ_0, θ_1 to reduce $J(\theta_0, \theta_1)$ until we hopefully end up at a minimum








Gradient descent algorithm

repeat until convergence {
 $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$ (for $j = 0$ and $j = 1$)
}

<pre>temp0 := $\theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$ temp1 := $\theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$ $\theta_0 :=$ temp0 $\theta_1 :=$ temp1</pre>	<pre>temp0 := $\theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$ $\theta_0 :=$ temp0 temp1 := $\theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$ $\theta_1 :=$ temp1</pre>
--	--

What, Why using a Derivative?

Why
?




Linear Regression: Quick Review

$$\theta_1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_1)$$

If α is too small, gradient descent can be slow.

If α is too large, gradient descent can overshoot the minimum. It may fail to converge, or even diverge.





Linear Regression: Quick Review

Gradient descent can converge to a local minimum, even with the learning rate α fixed.

$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$$

As we approach a local minimum, gradient descent will automatically take smaller steps. So, no need to decrease α over time.





Linear Regression: Quick Review

Gradient descent algorithm

repeat until convergence {
 $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$
 (for $j = 1$ and $j = 0$)
}

Linear Regression Model

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Linear Regression: Quick Review

$$\begin{aligned}\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) &= \frac{\partial}{\partial \theta_j} \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})^2 \\ &= \frac{\partial}{\partial \theta_j} \frac{1}{2m} \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2\end{aligned}$$

$$j = 0 : \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})$$

$$j = 1 : \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)}) x^{(i)}$$

Gradient descent algorithm

repeat until convergence {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$$

update
 θ_0 and θ_1
simultaneously

}

“Batch” Gradient Descent

“Batch”: Each step of gradient descent uses all the training examples.

Linear Regression

Introduction

Feature		Label	
	area	price	
	6.7	9.1	
	4.6	5.9	
	3.5	4.6	
	5.5	6.7	

House price data

Features			Label
TV	Radio	Newspaper	Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9

Advertising data

if area=6.0, price=?

if TV=55.0, Radio=34.0,
and Newspaper=62.0,
price=?

Features														Label
crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv	
0.00632	18	2.31	0	0.538	6.575	65.2	4.09	1	296	15.3	396.9	4.98	24	
0.02731	0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.9	9.14	21.6	
0.03237	0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4	
0.06905	0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.9	5.33	36.2	
0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	15.2	395.6	12.43	22.9	

Boston House Price Data

Linear Regression

❖ Area-based house price prediction

$\text{predicted_price} = w * \text{area} + b$


$\text{error} = (\text{predicted_price} - \text{real_price})^2$

$\hat{y}_i = wx_i + b$

$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$

$w = -0.34$
 $b = 0.04$

area	price	predicted	error
6.7	9.1	-2.238	128.55
4.6	5.9	-1.524	55.11
3.5	4.6	-1.15	33.06
5.5	6.7	-1.83	72.76



❖ Area-based house price prediction


predicted_price = w * area + b

error = (predicted_price – real_price)²

$$\hat{y}_i = wx_i + b$$

$$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

area	price	predicted	error
6.7	9.1	8.099	1.002
4.6	5.9	5.642	0.066
3.5	4.6	4.355	0.06
5.5	6.7	6.695	0.00002





Linear Regression

❖ Area-based house price prediction

$$\hat{y}_i = wx_i + b$$

$$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

How to change w and b so that $L(\hat{y}_i, y_i)$ reduces



Linear Regression

Linear equation

$$\hat{y} = wx + b$$

where \hat{y} is a predicted value,

w and b are parameters

and x is input feature

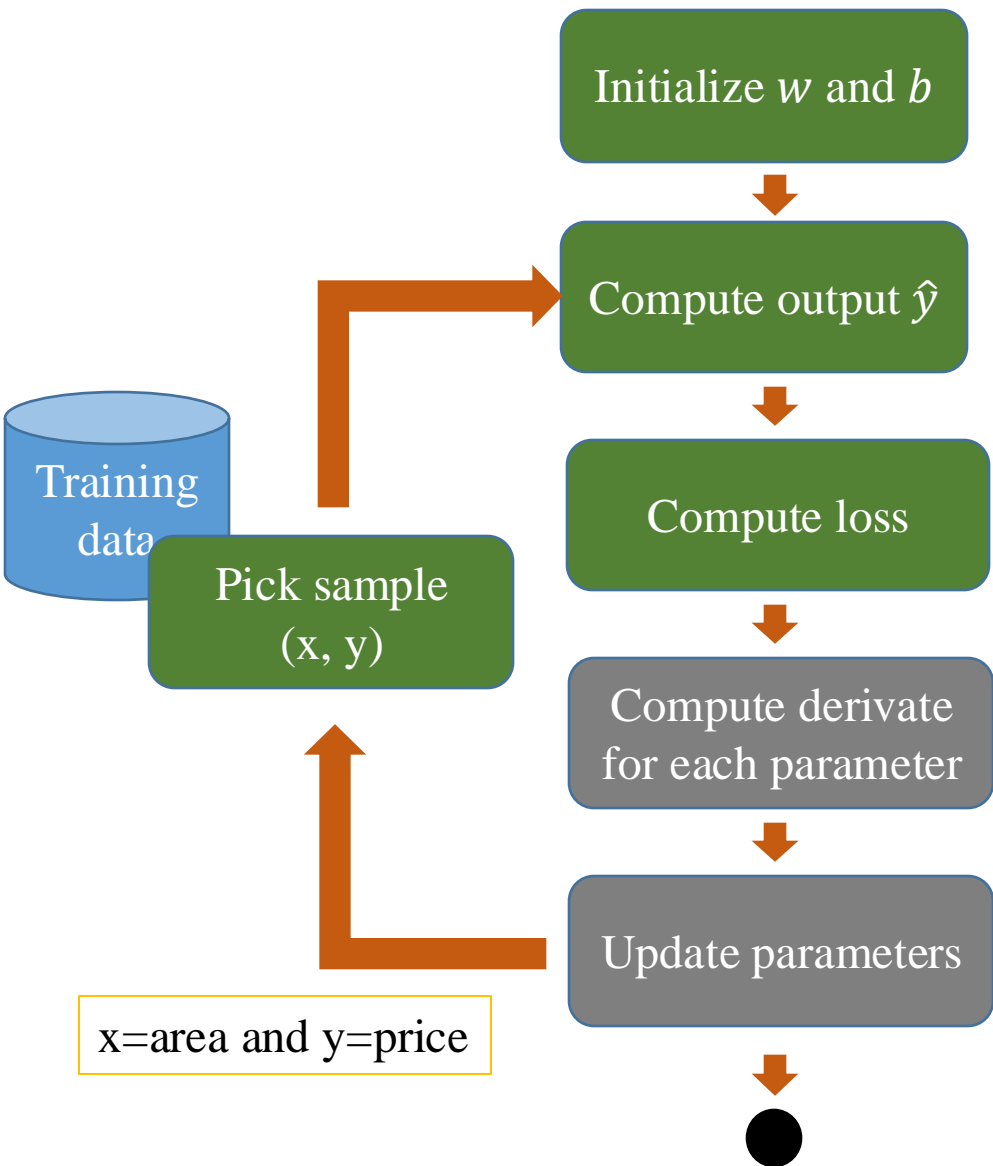
Error (loss) computation

Idea: compare predicted values \hat{y} and label values y

Squared loss

$$L(\hat{y}, y) = (\hat{y} - y)^2$$

How to find optimal w and b ?



Linear Regression

Linear equation

$$\hat{y} = wx + b$$

where \hat{y} is a predicted value,

w and b are parameters

and x is input feature

Error (loss) computation

Idea: compare predicted values \hat{y} and label values y

Squared loss

$$L(\hat{y}, y) = (\hat{y} - y)^2$$

How to find optimal w and b ?

Use gradient descent to minimize the loss function

Compute derivate for each parameter

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial b} = 2(\hat{y} - y)$$

Update parameters

$$w = w - \eta L'_w$$

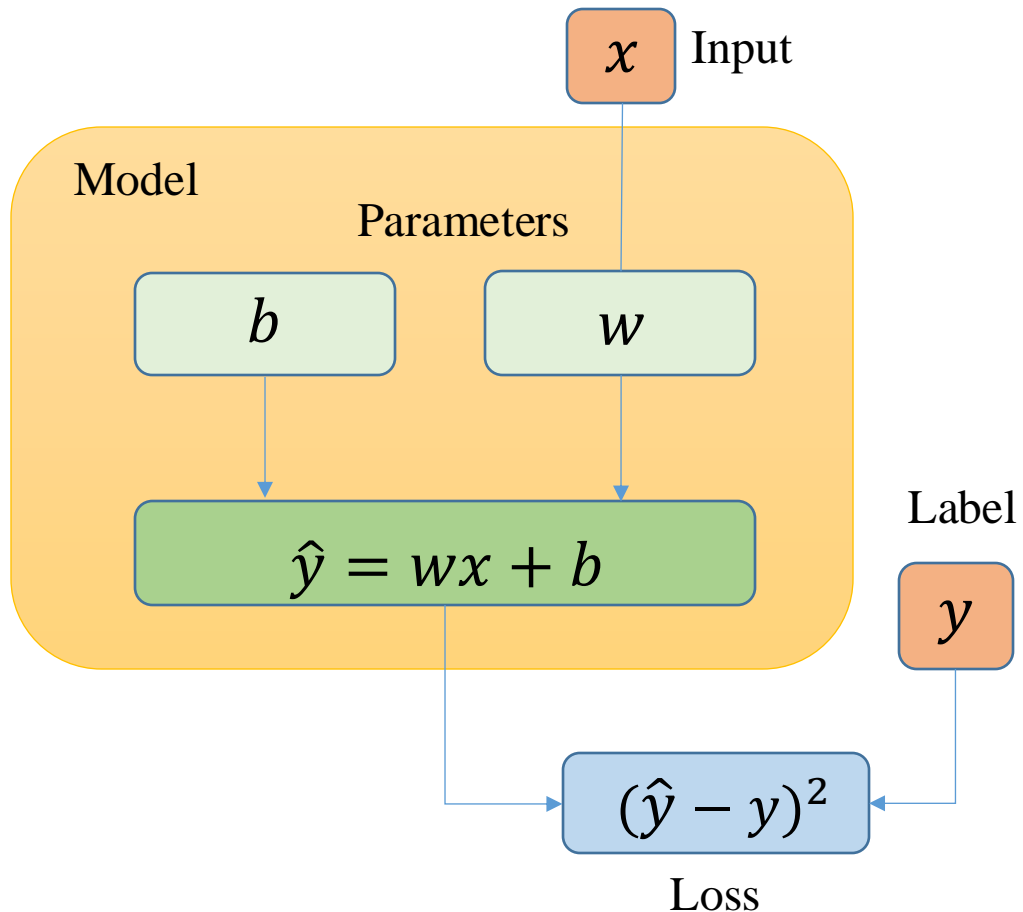
$$b = b - \eta L'_b$$

η is learning rate

Linear Regression

❖ Toy example

Diagram



Cheat sheet

Compute the output \hat{y}

$$\hat{y} = wx + b$$

Compute the loss

$$L = (\hat{y} - y)^2$$

Compute derivative

$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

Update parameters


$$w = w - \eta \frac{\partial L}{\partial w}$$

$$b = b - \eta \frac{\partial L}{\partial b}$$

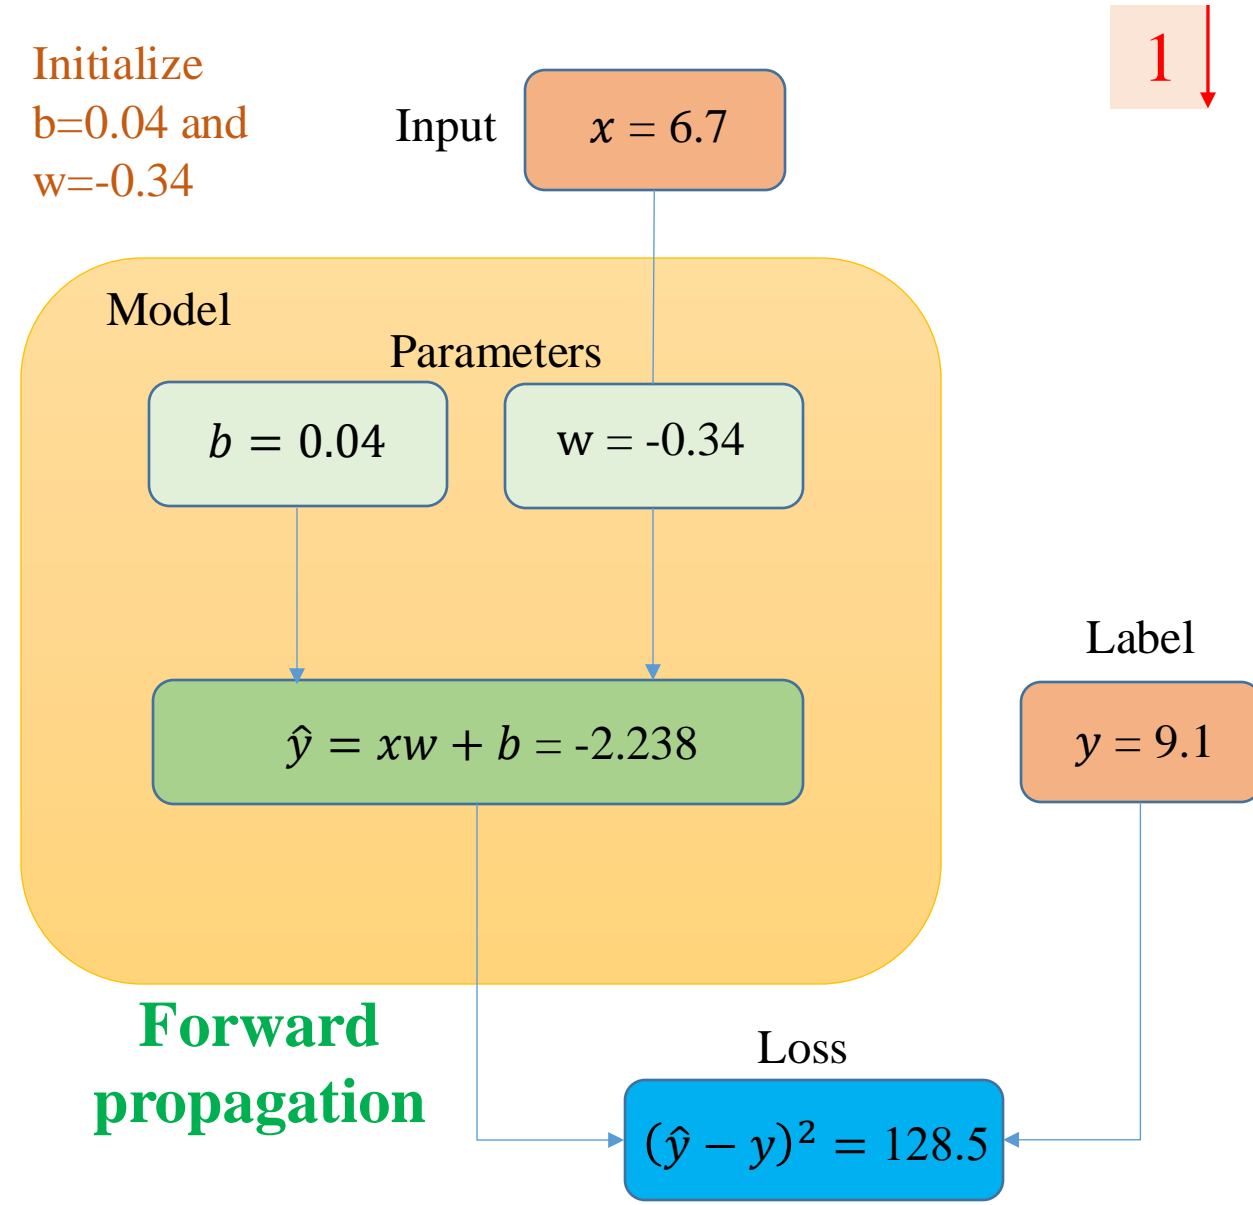
Linear Regression

Given
sample
data

Feature	Label
area	price
6.7	9.1
4.6	5.9
3.5	4.6
5.5	6.7

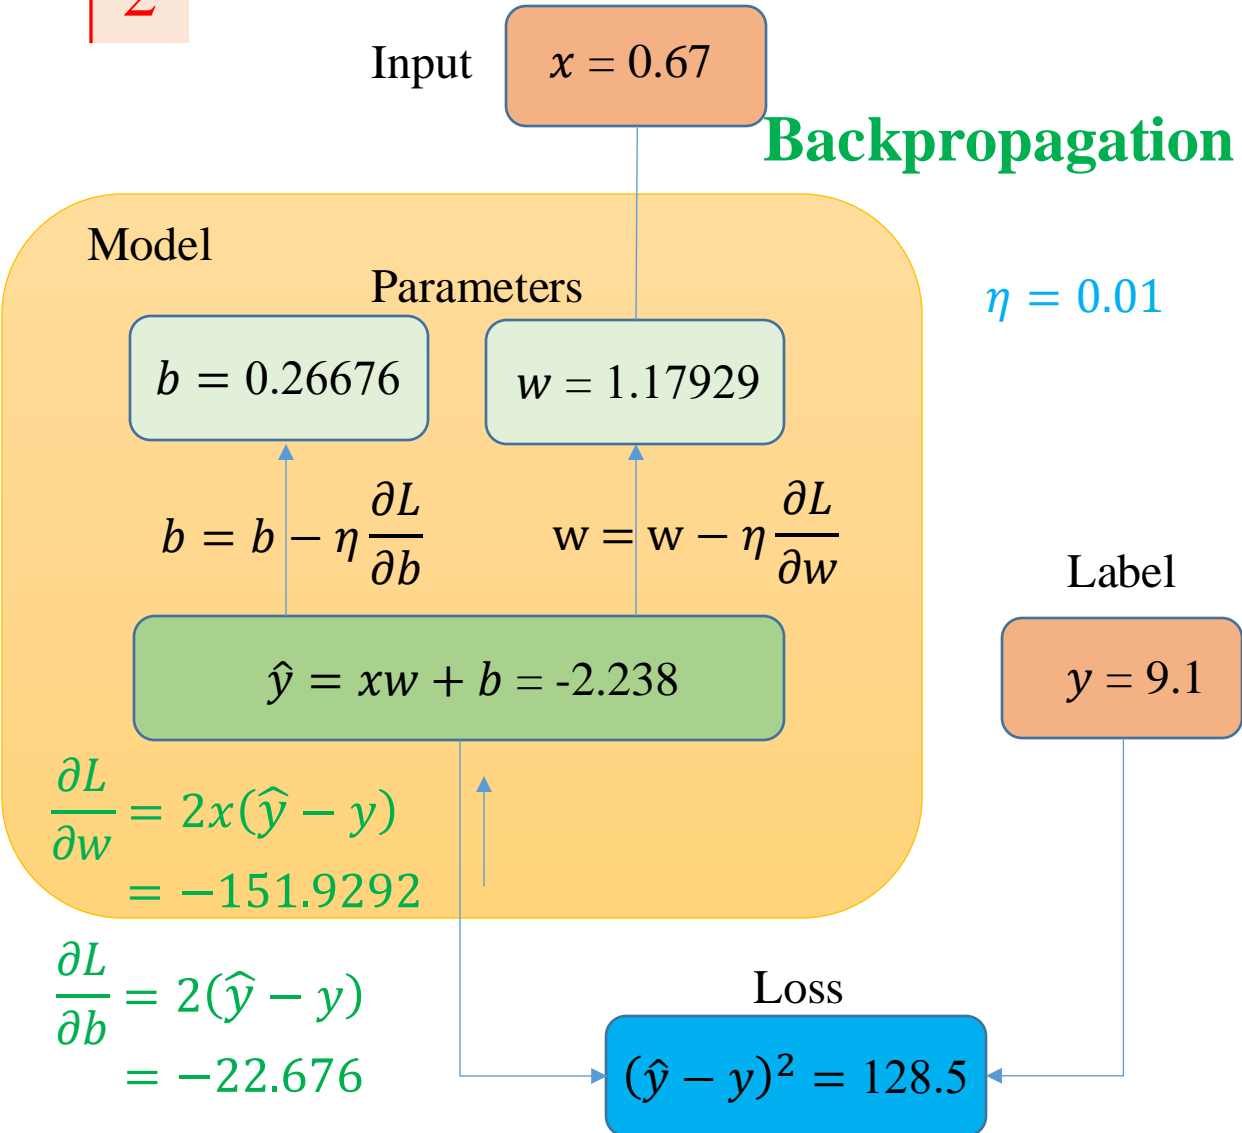


Initialize
 $b=0.04$ and
 $w=-0.34$

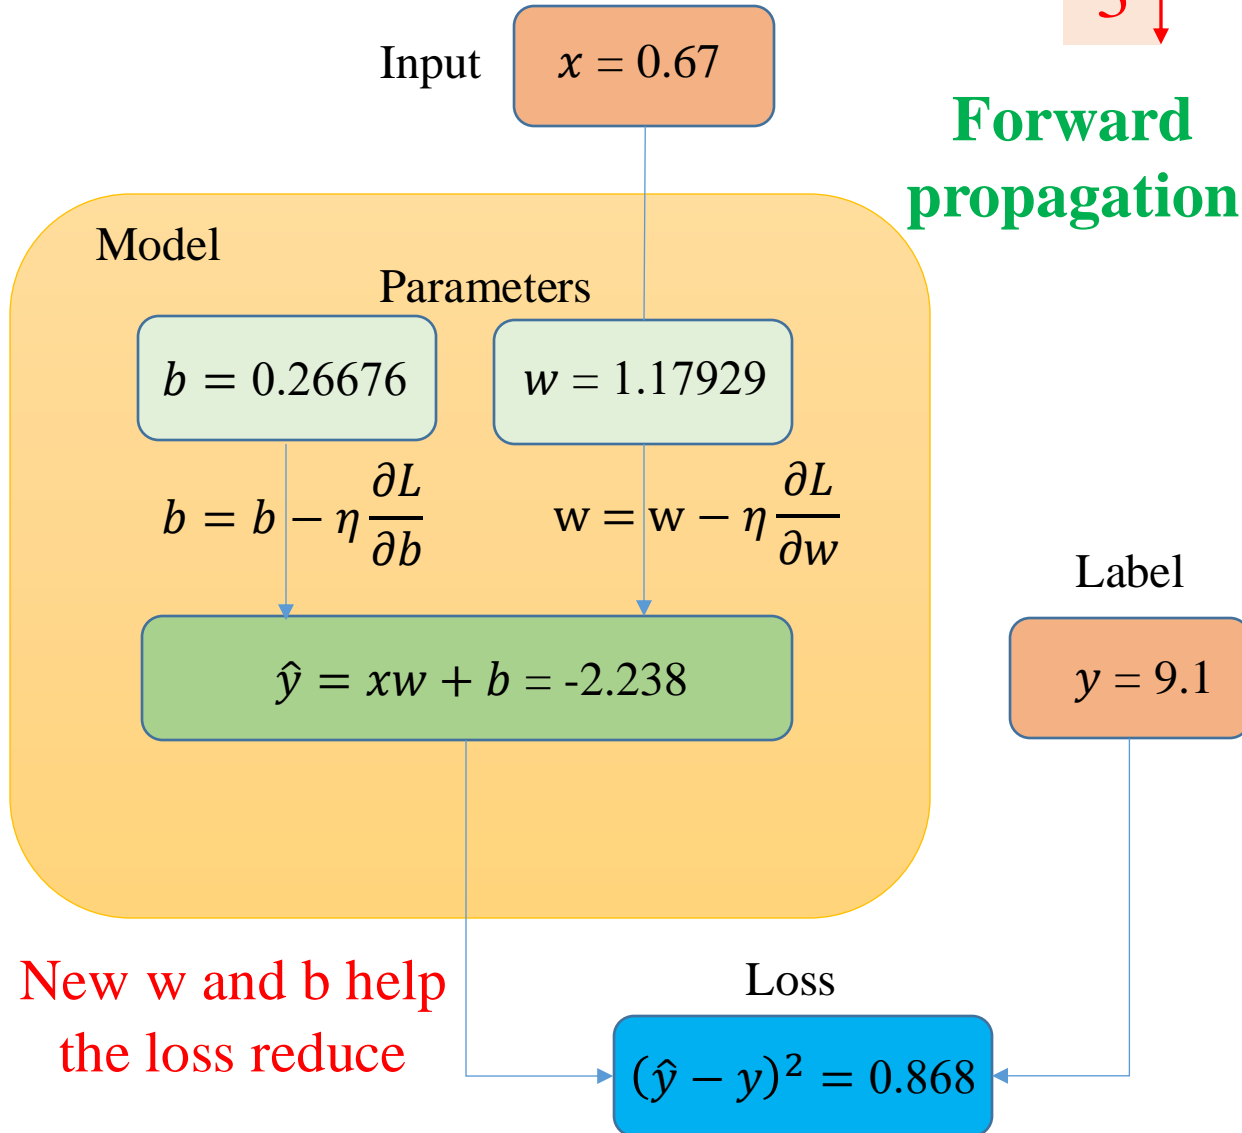


Linear Regression

2




3



Linear Regression


❖ Toy example

Model prediction before and after the first update



$$w = -0.34 \quad b = 0.04 \quad L = 128.55$$

Before updating

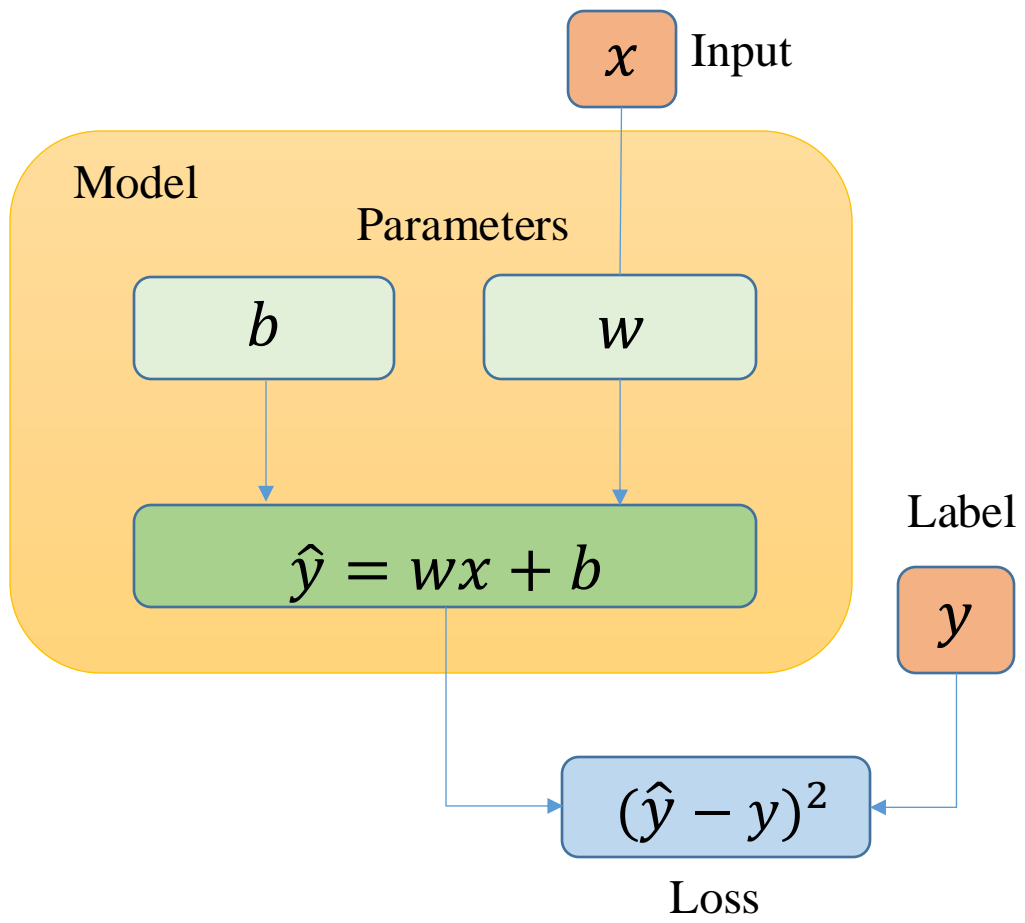


$$w = 1.179292 \quad b = 0.26676 \quad L = 0.868$$

After updating

Linear Regression

❖ Summary (simple version)



1) Pick a sample (x, y) from training data

2) Compute the output \hat{y}

$$\hat{y} = wx + b$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

5) Update parameters

$$w = w - \eta \frac{\partial L}{\partial w}$$

$$b = b - \eta \frac{\partial L}{\partial b}$$

η is learning rate

Linear Regression

❖ For the toy example

Cheat sheet

Compute the output \hat{y}

$$\hat{y} = wx + b$$

Compute the loss

$$L = (\hat{y} - y)^2$$

Compute derivative

$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

Update parameters


$$w = w - \eta \frac{\partial L}{\partial w}$$

$$b = b - \eta \frac{\partial L}{\partial b}$$

```

1  # forward
2  def predict(x, w, b):
3      return x*w + b
4
5  # compute gradient
6  def gradient(y_hat, y, x):
7      dw = 2*x*(y_hat-y)
8      db = 2*(y_hat-y)
9
10     return (dw, db)
11
12 # update weights
13 def update_weight(w, b, lr, dw, db):
14     w_new = w - lr*dw
15     b_new = b - lr*db
16
17     return (w_new, b_new)

```



Linear Regression

❖ Code for one update

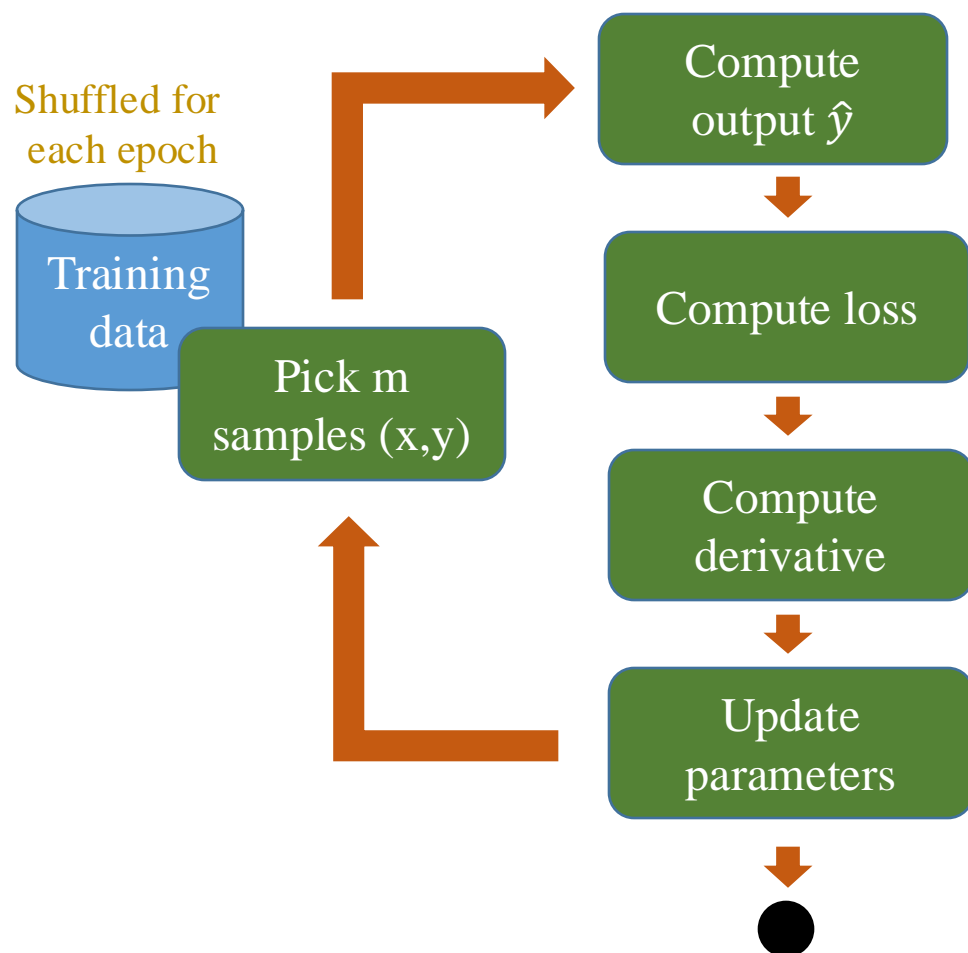
```
1 # forward
2 def predict(x, w, b):
3     return x*w + b
4
5 # compute gradient
6 def gradient(y_hat, y, x):
7     dw = 2*x*(y_hat-y)
8     db = 2*(y_hat-y)
9
10    return (dw, db)
11
12 # update weights
13 def update_weight(w, b, lr, dw, db):
14     w_new = w - lr*dw
15     b_new = b - lr*db
16
17    return (w_new, b_new)
```

```
1 # test sample
2 x = 6.7
3 y = 9.1
4
5 # init weights
6 b = 0.04
7 w = -0.34
8 lr = 0.01
9
10 # predict y_hat
11 y_hat = predict(x, w, b)
12 print('y_hat: ', y_hat)
13
14 # compute loss
15 loss = (y_hat-y)*(y_hat-y)
16 print('Loss: ', loss)
17
18 # compute gradient
19 (dw, db) = gradient(y_hat, y, x)
20 print('dw: ', dw)
21 print('db: ', db)
22
23 # update weights
24 (w, b) = update_weight(w, b, lr, dw, db)
25 print('w_new: ', w)
26 print('b_new: ', b)
```


Computational graph

❖ House price prediction

❖ m-sample training ($1 < m < N$)



1) Pick m samples $(x^{(i)}, y^{(i)})$ from training data

2) Tính output $\hat{y}^{(i)}$

$$\hat{y}^{(i)} = wx^{(i)} + b \quad \text{for } 0 \leq i < m$$

3) Tính loss

$$L^{(i)} = (\hat{y}^{(i)} - y^{(i)})^2 \quad \text{for } 0 \leq i < m$$

4) Tính đạo hàm

$$\frac{\partial L^{(i)}}{\partial w} = 2x^{(i)}(\hat{y}^{(i)} - y^{(i)})$$

$$\frac{\partial L^{(i)}}{\partial b} = 2(\hat{y}^{(i)} - y^{(i)})$$

for $0 \leq i < m$

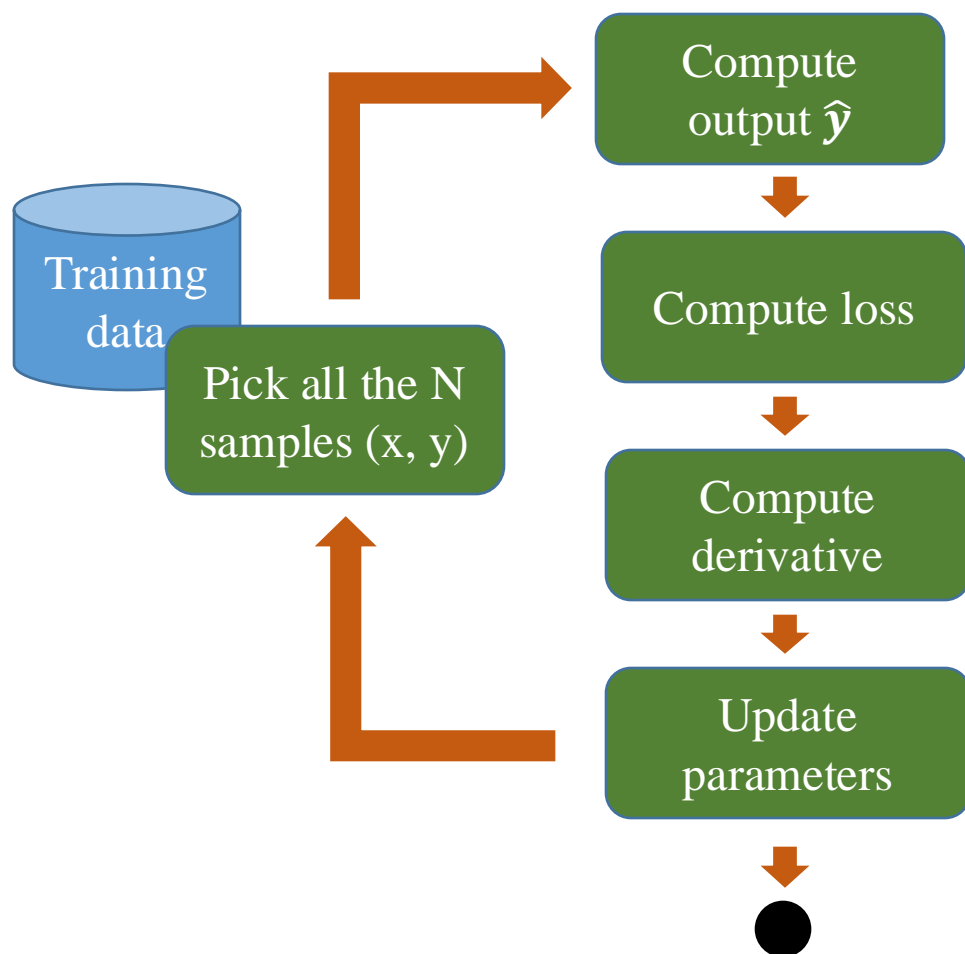
5) Cập nhật tham số

$$w = w - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial w}}{m} \quad b = b - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial b}}{m}$$

Computational graph

❖ House price prediction

❖ N-sample training



1) Pick all the N samples $(x^{(i)}, y^{(i)})$ from training data

2) Tính output $\hat{y}^{(i)}$

$$\hat{y}^{(i)} = wx^{(i)} + b \quad \text{for } 0 \leq i < N$$

3) Tính loss

$$L^{(i)} = (\hat{y}^{(i)} - y^{(i)})^2 \quad \text{for } 0 \leq i < N$$

4) Tính đạo hàm

$$\frac{\partial L^{(i)}}{\partial w} = 2x^{(i)}(\hat{y}^{(i)} - y^{(i)})$$

$$\frac{\partial L^{(i)}}{\partial b} = 2(\hat{y}^{(i)} - y^{(i)}) \quad \text{for } 0 \leq i < N$$

5) Cập nhật tham số

$$w = w - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial w}}{N} \quad b = b - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial b}}{N}$$

Linear Regression

❖ Generalized formula

House
price data

	Feature	Label	
	area	price	
	6.7	9.1	
	4.6	5.9	
	3.5	4.6	
	5.5	6.7	

Model

$$\text{price} = w * \text{area} + b$$

$$\hat{y} = wx + b$$

Model (vectorization)

$$\hat{y} = \boldsymbol{\theta}^T \mathbf{x} \quad \text{where} \quad \boldsymbol{\theta}^T = [b \quad w]^T$$
$$\mathbf{x} = [x_0 \quad \text{area}]^T$$
$$x_0 = 1$$

Advertising
data

Features			Label
TV	↕ Radio	↕ Newspaper	↕ Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9

Model

$$\text{Sale} = w_1 * TV + w_2 * Radio + w_3 * Newspaper + b$$

$$\hat{y} = w_1x_1 + w_2x_2 + w_3x_3 + b$$

Model (vectorization)

$$\hat{y} = \boldsymbol{\theta}^T \mathbf{x} \quad \text{where} \quad \boldsymbol{\theta}^T = [b \quad w_1 \quad w_2 \quad w_3]^T$$
$$\mathbf{x} = [x_0 \quad TV \quad Radio \quad Newspaper]^T$$
$$x_0 = 1$$

Linear Regression

1) Pick a sample (x, y) from training data

2) Compute the output \hat{y}

$$\hat{y} = w_1 * TV + w_2 * R + w_3 * N + b$$

$$\hat{y} = w_1 * x_1 + w_2 * x_2 + w_3 * x_3 + b$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\frac{\partial L}{\partial w_1} = 2x_1(\hat{y} - y) \quad \frac{\partial L}{\partial w_3} = 2x_3(\hat{y} - y)$$

$$\frac{\partial L}{\partial w_2} = 2x_2(\hat{y} - y) \quad \frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

5) Update parameters

$$w_1 = w_1 - \eta \frac{\partial L}{\partial w_1} \quad w_3 = w_3 - \eta \frac{\partial L}{\partial w_3}$$

$$w_2 = w_2 - \eta \frac{\partial L}{\partial w_2} \quad b = b - \eta \frac{\partial L}{\partial b}$$


Features			Label
TV	↕ Radio	↕ Newspaper	↕ Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9

Advertising
data

Model

$$\text{Sale} = w_1 * TV + w_2 * Radio + w_3 * Newspaper + b$$

$$\hat{y} = w_1x_1 + w_2x_2 + w_3x_3 + b$$



Linear Regression

Features			Label
TV	Radio	Newspaper	Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9

```

1 def initialize_params():
2     w1 = random.gauss(mu=0.0, sigma=0.01)
3     w2 = random.gauss(mu=0.0, sigma=0.01)
4     w3 = random.gauss(mu=0.0, sigma=0.01)
5     b = 0
6
7     return w1, w2, w3, b
8
9 # initialize model's parameters
10 w1, w2, w3, b = initialize_params()
  
```

```

1 # compute output and loss
2 def predict(x1, x2, x3, w1, w2, w3, b):
3     return w1*x1 + w2*x2 + w3*x3 + b
4
5 def compute_loss(y_hat, y):
6     return (y_hat - y)**2
7
8 # compute gradient
9 def compute_gradient_wi(xi, y, y_hat):
10    dl_dwi = 2*xi*(y_hat-y)
11    return dl_dwi
12
13 def compute_gradient_b(y, y_hat):
14    dl_db = 2*(y_hat-y)
15    return dl_db
16
17 # update weights
18 def update_weight_wi(wi, dl_dwi, lr):
19    wi = wi - lr*dl_dwi
20    return wi
21
22 def update_weight_b(b, dl_db, lr):
23    b = b - lr*dl_db
24    return b
  
```

Outline



➤ **Linear Regression Review**

➤ **Exercise 1**

➤ **Exercise 2**

➤ **Exercise 3**

➤ **Exercise 4**

➤ **Exercise 5**

➤ **Other Discussions**

Exercise 1

Bài tập 1 (kỹ thuật đọc và xử lý dữ liệu từ file .csv): Cho trước file dữ liệu advertising.csv, hãy hoàn thành function **prepare_data(file_name_dataset)** trả về dữ liệu đã được tổ chức (X cho input và y cho output).

```
1
2 # dataset
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import random
6
7 def get_column(data, index):
8
9     #your code here *****
10
11     return result
12
13 def prepare_data(file_name_dataset):
14     data = np.genfromtxt(file_name_dataset, delimiter=',', skip_header=1).tolist()
15     N = len(data)
16
17     # get tv (index=0)
18     tv_data = get_column(data, 0)
19
20     # get radio (index=1)
21     radio_data = get_column(data, 1)
22
```

Exercise 1

TV	Radio	Newspaper	Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9
8.7	48.9	75	7.2
57.5	32.8	23.5	11.8
120.2	19.6	11.6	13.2
8.6	2.1	1	4.8
199.8	2.6	21.2	15.6

X

Y

TV= X[0] = [230.1, 44.5, 17.2, 151.5, 180.8, 8.7, 57.5, 120.2, 8.6, 199.8]

Radio = X[1] = [37.8, 39.9, 45.9, 41.3, 10.8, 48.9, 32.8, 19.6, 2,1, 2.6]

News = X[2] = [69.2, 45.1, 69.3, 58.5, 58.4, 75, 32.5, 11.6, 1. 0, 21.2]

Sales = Y = [22.1, 10.4, 12, 16.5, 17.9, 7.2, 11.8, 13.2, 4.8, 15.6]

Exercise 1

```
▶ # dataset
import numpy as np
import matplotlib.pyplot as plt
import random

def get_column(data, index):
    result = [row[index] for row in data]
    return result

def prepare_data(file_name_dataset):
    data = np.genfromtxt(file_name_dataset, delimiter=',', skip_header=1).tolist()
    N = len(data)

    # get tv (index=0)
    tv_data = get_column(data, 0)

    # get radio (index=1)
    radio_data = get_column(data, 1)

    # get newspaper (index=2)
    newspaper_data = get_column(data, 2)

    # get sales (index=3)
    sales_data = get_column(data, 3)

    # building X input and y output for training
    X = [tv_data, radio_data, newspaper_data]
    y = sales_data
    return X,y
```

Outline

➤ **Linear Regression Review**

➤ **Exercise 1**

➤ **Exercise 2**

➤ **Exercise 3**

➤ **Exercise 4**

➤ **Exercise 5**

➤ **Other Discussions**



Exercise 2

Bài tập 2 (kỹ thuật huấn luyện data dùng one sample - linear regression): Sử dụng kết quả dữ liệu đầu vào X , và dữ liệu đầu ra y từ bài 1, để phát triển chương trình dự đoán thông tin sales (y) từ X bằng cách dùng giải thuật linear regression with one sample-training với loss được tính bằng công thức Mean Squared Error $L = (\hat{y} - y)^2$. Sơ đồ hoạt động của giải thuật được mô tả ở hình 2. Nhiệm vụ của bạn là hoàn thành function `implement_linear_regression(X_data, y_data, epoch_max, lr)` và trả về 4 tham số w_1, w_2, w_3, b và lịch sử tính loss như bên dưới.

1) Pick a sample (x, y) from training data

2) Tính output \hat{y}

$$\hat{y} = w_1 x_1 + w_2 x_2 + w_3 x_3 + b$$

3) Tính loss

$$L = (\hat{y} - y)^2$$

4) Tính đạo hàm

$$L'_{w_1} = 2x_1(\hat{y} - y)$$

$$L'_{w_2} = 2x_2(\hat{y} - y)$$

$$L'_{w_3} = 2x_3(\hat{y} - y)$$

$$L'_b = 2(\hat{y} - y)$$

5) Cập nhật tham số

$$w_1 = w_1 - \eta L'_{w_1}$$

$$w_2 = w_2 - \eta L'_{w_2}$$

$$w_3 = w_3 - \eta L'_{w_3}$$

$$b = b - \eta L'_b$$

η is learning rate

Hình 2: Các bước để thực hiện train linear regression model



Exercise 2

```


1 def implement_linear_regression(X_data, y_data, epoch_max = 50, lr = 1e-5):
2     losses = []
3
4     w1, w2, w3, b = initialize_params()
5
6     N = len(y_data)
7     for epoch in range(epoch_max):
8         for i in range(N):
9             # get a sample
10            x1 = X_data[0][i]
11            x2 = X_data[1][i]
12            x3 = X_data[2][i]
13
14            y = y_data[i]
15
16            # print(y)
17            # compute output
18            y_hat = predict(x1, x2, x3, w1, w2, w3, b)
19
20            # compute loss
21            loss = compute_loss_mse(y, y_hat)
22
23            # compute gradient w1, w2, w3, b
24            dl_dw1 = compute_gradient_wi(x1, y, y_hat)
25            dl_dw2 = compute_gradient_wi(x2, y, y_hat)
26            dl_dw3 = compute_gradient_wi(x3, y, y_hat)
27            dl_db = compute_gradient_b(y, y_hat)
28
29            # update parameters
30            w1 = update_weight_wi(w1, dl_dw1, lr)
31            w2 = update_weight_wi(w2, dl_dw2, lr)
32            w3 = update_weight_wi(w3, dl_dw3, lr)
33            b = update_weight_b(b, dl_db, lr)
34
35            # logging
36            losses.append(loss)
37
38    return (w1, w2, w3, b, losses)

```

```

1 def initialize_params():
2     # w1 = random.gauss(mu=0.0, sigma=0.01)
3     # w2 = random.gauss(mu=0.0, sigma=0.01)
4     # w3 = random.gauss(mu=0.0, sigma=0.01)
5     # b = 0
6
7     w1, w2, w3, b = (0.016992259082509283, 0.0070783670518262355,
8                     -0.002307860847821344, 0)
9
10    return w1, w2, w3, b

```




Exercise 2

```
1 def implement_linear_regression(X_data, y_data, epoch_max = 50, lr = 1e-5):
2     losses = []
3
4     w1, w2, w3, b = initialize_params()
5
6     N = len(y_data)
7     for epoch in range(epoch_max):
8         for i in range(N):
9             # get a sample
10            x1 = X_data[0][i]
11            x2 = X_data[1][i]
12            x3 = X_data[2][i]
13
14            y = y_data[i]
15
16            # print(y)
17            # compute output
18            y_hat = predict(x1, x2, x3, w1, w2, w3, b)
19
20            # compute loss
21            loss = compute_loss_mse(y, y_hat)
22
23            # compute gradient w1, w2, w3, b
24            dl_dw1 = compute_gradient_wi(x1, y, y_hat)
25            dl_dw2 = compute_gradient_wi(x2, y, y_hat)
26            dl_dw3 = compute_gradient_wi(x3, y, y_hat)
27            dl_db = compute_gradient_b(y, y_hat)
28
29            # update parameters
30            w1 = update_weight_wi(w1, dl_dw1, lr)
31            w2 = update_weight_wi(w2, dl_dw2, lr)
32            w3 = update_weight_wi(w3, dl_dw3, lr)
33            b = update_weight_b(b, dl_db, lr)
34
35            # logging
36            losses.append(loss)
37    return (w1, w2, w3, b, losses)
```



compute output and loss


```
def predict(x1, x2, x3, w1, w2, w3, b):
    return w1*x1 + w2*x2 + w3*x3 + b
```

Exercise 2

```
1 def implement_linear_regression(X_data, y_data, epoch_max = 50, lr = 1e-5):
2     losses = []
3
4     w1, w2, w3, b = initialize_params()
5
6     N = len(y_data)
7     for epoch in range(epoch_max):
8         for i in range(N):
9             # get a sample
10            x1 = X_data[0][i]
11            x2 = X_data[1][i]
12            x3 = X_data[2][i]
13
14            y = y_data[i]
15
16            # print(y)
17            # compute output
18            y_hat = predict(x1, x2, x3, w1, w2, w3, b)
19
20            # compute loss
21            loss = compute_loss_mse(y, y_hat)
22
23            # compute gradient w1, w2, w3, b
24            dl_dw1 = compute_gradient_wi(x1, y, y_hat)
25            dl_dw2 = compute_gradient_wi(x2, y, y_hat)
26            dl_dw3 = compute_gradient_wi(x3, y, y_hat)
27            dl_db = compute_gradient_b(y, y_hat)
28
29            # update parameters
30            w1 = update_weight_wi(w1, dl_dw1, lr)
31            w2 = update_weight_wi(w2, dl_dw2, lr)
32            w3 = update_weight_wi(w3, dl_dw3, lr)
33            b = update_weight_b(b, dl_db, lr)
34
35            # logging
36            losses.append(loss)
37
38 return (w1, w2, w3, b, losses)
```

```
def compute_loss_mse(y_hat, y):
    return (y_hat - y)**2
```



Exercise 2

```
1 def implement_linear_regression(X_data, y_data, epoch_max = 50, lr = 1e-5):
2     losses = []
3
4     w1, w2, w3, b = initialize_params()
5
6     N = len(y_data)
7     for epoch in range(epoch_max):
8         for i in range(N):
9             # get a sample
10            x1 = X_data[0][i]
11            x2 = X_data[1][i]
12            x3 = X_data[2][i]
13
14            y = y_data[i]
15
16            # print(y)
17            # compute output
18            y_hat = predict(x1, x2, x3, w1, w2, w3, b)
19
20            # compute loss
21            loss = compute_loss_mse(y, y_hat)
22
23            # compute gradient w1, w2, w3, b
24            dl_dw1 = compute_gradient_wi(x1, y, y_hat)
25            dl_dw2 = compute_gradient_wi(x2, y, y_hat)
26            dl_dw3 = compute_gradient_wi(x3, y, y_hat)
27            dl_db = compute_gradient_b(y, y_hat)
28
29            # update parameters
30            w1 = update_weight_wi(w1, dl_dw1, lr)
31            w2 = update_weight_wi(w2, dl_dw2, lr)
32            w3 = update_weight_wi(w3, dl_dw3, lr)
33            b = update_weight_b(b, dl_db, lr)
34
35            # logging
36            losses.append(loss)
37    return (w1, w2, w3, b, losses)
```


```
# compute gradient
def compute_gradient_wi(xi, y, y_hat):
    dl_dwi = 2*xi*(y_hat-y)
    return dl_dwi
```



Exercise 2

```
1 def implement_linear_regression(X_data, y_data, epoch_max = 50, lr = 1e-5):
2     losses = []
3
4     w1, w2, w3, b = initialize_params()
5
6     N = len(y_data)
7     for epoch in range(epoch_max):
8         for i in range(N):
9             # get a sample
10            x1 = X_data[0][i]
11            x2 = X_data[1][i]
12            x3 = X_data[2][i]
13
14            y = y_data[i]
15
16            # print(y)
17            # compute output
18            y_hat = predict(x1, x2, x3, w1, w2, w3, b)
19
20            # compute loss
21            loss = compute_loss_mse(y, y_hat)
22
23            # compute gradient w1, w2, w3, b
24            dl_dw1 = compute_gradient_wi(x1, y, y_hat)
25            dl_dw2 = compute_gradient_wi(x2, y, y_hat)
26            dl_dw3 = compute_gradient_wi(x3, y, y_hat)
27            dl_db = compute_gradient_b(y, y_hat)
28
29            # update parameters
30            w1 = update_weight_wi(w1, dl_dw1, lr)
31            w2 = update_weight_wi(w2, dl_dw2, lr)
32            w3 = update_weight_wi(w3, dl_dw3, lr)
33            b = update_weight_b(b, dl_db, lr)
34
35            # logging
36            losses.append(loss)
37    return (w1, w2, w3, b, losses)
```

```
def compute_gradient_b(y, y_hat):
    dl_db = 2*(y_hat-y)
    return dl_db
```

Exercise 2

```
1 def implement_linear_regression(X_data, y_data, epoch_max = 50, lr = 1e-5):
2     losses = []
3
4     w1, w2, w3, b = initialize_params()
5
6     N = len(y_data)
7     for epoch in range(epoch_max):
8         for i in range(N):
9             # get a sample
10            x1 = X_data[0][i]
11            x2 = X_data[1][i]
12            x3 = X_data[2][i]
13
14            y = y_data[i]
15
16            # print(y)
17            # compute output
18            y_hat = predict(x1, x2, x3, w1, w2, w3, b)
19
20            # compute loss
21            loss = compute_loss_mse(y, y_hat)
22
23            # compute gradient w1, w2, w3, b
24            dl_dw1 = compute_gradient_wi(x1, y, y_hat)
25            dl_dw2 = compute_gradient_wi(x2, y, y_hat)
26            dl_dw3 = compute_gradient_wi(x3, y, y_hat)
27            dl_db = compute_gradient_b(y, y_hat)
28
29            # update parameters
30            w1 = update_weight_wi(w1, dl_dw1, lr)
31            w2 = update_weight_wi(w2, dl_dw2, lr)
32            w3 = update_weight_wi(w3, dl_dw3, lr)
33            b = update_weight_b(b, dl_db, lr)
34
35            # logging
36            losses.append(loss)
37    return (w1, w2, w3, b, losses)
```

```
# update weights
def update_weight_wi(wi, dl_dwi, lr):
    wi = wi - lr*dl_dwi
    return wi
```





Exercise 2

```
1 def implement_linear_regression(X_data, y_data, epoch_max = 50, lr = 1e-5):
2     losses = []
3
4     w1, w2, w3, b = initialize_params()
5
6     N = len(y_data)
7     for epoch in range(epoch_max):
8         for i in range(N):
9             # get a sample
10            x1 = X_data[0][i]
11            x2 = X_data[1][i]
12            x3 = X_data[2][i]
13
14            y = y_data[i]
15
16            # print(y)
17            # compute output
18            y_hat = predict(x1, x2, x3, w1, w2, w3, b)
19
20            # compute loss
21            loss = compute_loss_mse(y, y_hat)
22
23            # compute gradient w1, w2, w3, b
24            dl_dw1 = compute_gradient_wi(x1, y, y_hat)
25            dl_dw2 = compute_gradient_wi(x2, y, y_hat)
26            dl_dw3 = compute_gradient_wi(x3, y, y_hat)
27            dl_db = compute_gradient_b(y, y_hat)
28
29            # update parameters
30            w1 = update_weight_wi(w1, dl_dw1, lr)
31            w2 = update_weight_wi(w2, dl_dw2, lr)
32            w3 = update_weight_wi(w3, dl_dw3, lr)
33            b = update_weight_b(b, dl_db, lr)
34
35            # logging
36            losses.append(loss)
37
38     return (w1, w2, w3, b, losses)
```

```
def update_weight_b(b, dl_db, lr):
    b = b - lr*dl_db
    return b
```

Exercise 2





Exercise 2

```
# given new data
tv = 19.2
radio = 35.9
newspaper = 51.3

X,y = prepare_data('advertising.csv')
(w1,w2,w3,b, losses) = implement_linear_regression(X,y)
sales = predict(tv, radio, newspaper, w1, w2, w3, b)
print(f'predicted sales is {sales}')
```

Outline

➤ **Linear Regression Review**

➤ **Exercise 1**

➤ **Exercise 2**

➤ **Exercise 3**

➤ **Exercise 4**

➤ **Exercise 5**

➤ **Other Discussions**



Exercise 3


Bài tập 3 (kỹ thuật huấn luyện data dùng batch N samples - linear regression): Cải tiến giải thuật ở bài tập 2, bằng cách huấn luyện giải thuật linear regression sử dụng N samples-training. Công việc của bạn ở bài tập này là bạn cần implement lại function **implement_linear_regression_nsamples** sử dụng N sample-training với MSE loss function $L(\hat{y}, y) = \frac{1}{N} \sum_{i=1}^N (\hat{y} - y)^2$ và MAE loss function $L(\hat{y}, y) = \frac{1}{N} \sum_{i=1}^N |\hat{y} - y|$

```
1 def implement_linear_regression_nsamples(X_data, y_data, epoch_max = 50, lr = 1e-5):  
2     losses = []  
3  
4     w1, w2, w3, b = initialize_params()  
5     N = len(y_data)
```




Exercise 3

```
def implement_linear_regression_nsamples(X_data, y_data, epoch_max = 50, lr = 1e-5):  
    losses = []  
  
    w1, w2, w3, b = initialize_params()  
    N = len(y_data)  
  
    for epoch in range(epoch_max):  
  
        loss_total = 0.0  
        dw1_total = 0.0  
        dw2_total = 0.0  
        dw3_total = 0.0  
        db_total = 0.0  
  
        for i in range(N):  
            # get a sample  
            x1 = X_data[0][i]  
            x2 = X_data[1][i]  
            x3 = X_data[2][i]  
  
            y = y_data[i]  
  
            # print(y)  
            # compute output  
            y_hat = predict(x1, x2, x3, w1, w2, w3, b)
```



Exercise 3

```
for i in range(N):
    # get a sample
    x1 = X_data[0][i]
    x2 = X_data[1][i]
    x3 = X_data[2][i]

    y = y_data[i]

    # print(y)
    # compute output
    y_hat = predict(x1, x2, x3, w1, w2, w3, b)

    # compute loss
    loss = compute_loss_mae(y, y_hat)
    loss_total = loss_total + loss

    # compute gradient w1, w2, w3, b
    dl_dw1 = compute_gradient_wi(x1, y, y_hat)
    dl_dw2 = compute_gradient_wi(x2, y, y_hat)
    dl_dw3 = compute_gradient_wi(x3, y, y_hat)
    dl_db = compute_gradient_b(y, y_hat)


    # accumulate
    dw1_total = dw1_total + dl_dw1
    dw2_total = dw2_total + dl_dw2
    dw3_total = dw3_total + dl_dw3
    db_total = db_total + dl_db
```

```
# (after processing N samples) - update parameters
w1 = update_weight_wi(w1, dl_dw1/N, lr)
w2 = update_weight_wi(w2, dl_dw2/N, lr)
w3 = update_weight_wi(w3, dl_dw3/N, lr)
b = update_weight_b(b, dl_db/N, lr)


# logging
losses.append(loss_total/N)
return (w1, w2, w3, b, losses)
```


Exercise 3

```
X,y = prepare_data('advertising.csv')
(w1,w2,w3,b, losses) = implement_linear_regression_nsamples(X,y,1000)
plt.plot(losses)
plt.xlabel("#epoch")
plt.ylabel("MAE Loss")
plt.show()
```



```
X,y = prepare_data('advertising.csv')
(w1,w2,w3,b, losses) = implement_linear_regression_nsamples(X,y,1000)
plt.plot(losses)
plt.xlabel("#epoch")
plt.ylabel("MSE Loss")
plt.show()
```



Outline

➤ **Linear Regression Review**

➤ **Exercise 1**

➤ **Exercise 2**

➤ **Exercise 3**

➤ **Exercise 4**

➤ **Exercise 5**

➤ **Other Discussions**



Exercise 4

Bài tập 4 Như chúng ta đã biết, mục đích của linear regression là tìm hàm xấp xỉ $y = ax_1 + bx_2 + cx_3 + bx_0$. Trong đó x_1 là TV, x_2 là Radio, x_3 là Newspapers, và $x_0 = 1$. Đầu tiên, bạn cần tổ chức lại dữ liệu đầu vào ở bài tập 1 theo dạng danh sách các feature (x_0, x_1, x_2, x_3). Ví dụ theo hình 1, dữ liệu đầu vào dòng thứ 1 và 2 ta có thể tổ chức lại như sau:

$$X[0] = [1, x_1, x_2, x_3] = [1, 230.1, 37.8, 69.2]$$

$$X[1] = [1, x_1, x_2, x_3] = [1, 44.5, 39.3, 45.1]$$

10

AI VIETNAM

aivietnam.edu.vn

....

$$X[199] = [1, x_1, x_2, x_3] = [1, 232.1, 8.6, 8.7]$$

Để implement ý tưởng trên vào chương trình, bạn có thể sử dụng function bên dưới:

Linear Regression

❖ Generalized formula

House
price data

	Feature	Label	
	area	price	
	6.7	9.1	
	4.6	5.9	
	3.5	4.6	
	5.5	6.7	

Model

$$\text{price} = w * \text{area} + b$$

$$\hat{y} = wx + b$$

Model (vectorization)

$$\hat{y} = \theta^T x \quad \text{where} \quad \theta^T = [b \quad w]^T$$
$$x = [x_0 \quad \text{area}]^T$$
$$x_0 = 1$$

Advertising
data

Features			Label
TV	Radio	Newspaper	Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9

Model

$$\text{Sale} = w_1 * TV + w_2 * Radio + w_3 * Newspaper + b$$

$$\hat{y} = w_1 x_1 + w_2 x_2 + w_3 x_3 + b$$

Model (vectorization)

$$\hat{y} = \theta^T x \quad \text{where} \quad \theta^T = [b \quad w_1 \quad w_2 \quad w_3]^T$$
$$x = [x_0 \quad TV \quad Radio \quad Newspaper]^T$$
$$x_0 = 1$$

Exercise 4

```
▶ def prepare_data(file_name_dataset):  
    data = np.genfromtxt(file_name_dataset, delimiter=',', skip_header=1).tolist()  
  
    # get tv (index=0)  
    tv_data = get_column(data, 0)  
  
    # get radio (index=1)  
    radio_data = get_column(data, 1)  
  
    # get newspaper (index=2)  
    newspaper_data = get_column(data, 2)  
  
    # get sales (index=3)  
    sales_data = get_column(data, 3)  
  
    # building X input and y output for training  
    # Create list of features for input  
    X = [[1, x1, x2, x3] for x1, x2, x3 in zip(tv_data, radio_data, newspaper_data)]  
    y = sales_data  
    return X, y
```

Exercise 4

```
def implement_linear_regression(X_feature, y_output, epoch_max = 50, lr = 1e-5):
```

```
    losses = []
```

```
    weights = initialize_params()
```

```
    N = len(y_output)
```

```
    for epoch in range(epoch_max):
```

```
        print("epoch", epoch)
```

```
        for i in range(N):
```

```
            # get a sample - row i
```

```
            features_i = X_feature[i]
```

```
            y = sales_data[i]
```

```
            # compute output
```

```
            y_hat = predict(features_i, weights)
```

```
            # compute loss
```

```
            loss = compute_loss(y, y_hat)
```

```
            # compute gradient w1, w2, w3, b
```

```
            dl_dweights = compute_gradient_w(features_i, y, y_hat)
```

```
            # update parameters
```

```
            weights = update_weight(weights, dl_dweights, lr)
```


```
            # logging
```

```
            losses.append(loss)
```

```
    return weights, losses
```

```
def predict(X_features, weights):
```

```
    return sum([f*w for f, w in zip(X_features, weights)])
```



Exercise 4

```
▶ def implement_linear_regression(X_feature, y_output, epoch_max = 50, lr = 1e-5):
```

```
    losses = []
```

```
    weights = initialize_params()
```

```
    N = len(y_output)
```

```
    for epoch in range(epoch_max):
```

```
        print("epoch", epoch)
```

```
        for i in range(N):
```

```
            # get a sample - row i
```

```
            features_i = X_feature[i]
```

```
            y = sales_data[i]
```

```
            # compute output
```

```
            y_hat = predict(features_i, weights)
```

```
            # compute loss
```

```
            loss = compute_loss(y, y_hat)
```

```
            # compute gradient w1, w2, w3, b
```

```
            dl_dweights = compute_gradient_w(features_i, y, y_hat)
```

```
            # update parameters
```

```
            weights = update_weight(weights, dl_dweights, lr)
```

```
            # logging
```

```
            losses.append(loss)
```

```
    return weights, losses
```

```
def compute_loss(y_hat, y):  
    return (y_hat - y)**2
```


Exercise 4

```
def implement_linear_regression(X_feature, y_output, epoch_max = 50, lr = 1e-5):
```

```
    losses = []
```

```
    weights = initialize_params()
```

```
    N = len(y_output)
```

```
    for epoch in range(epoch_max):
```

```
        print("epoch", epoch)
```

```
        for i in range(N):
```

```
            # get a sample - row i
```

```
            features_i = X_feature[i]
```

```
            y = sales_data[i]
```

```
            # compute output
```

```
            y_hat = predict(features_i, weights)
```

```
            # compute loss
```

```
            loss = compute_loss(y, y_hat)
```

```
            # compute gradient w1, w2, w3, b
```

```
            dl_dweights = compute_gradient_w(features_i, y, y_hat)
```

```
            # update parameters
```

```
            weights = update_weight(weights, dl_dweights, lr)
```

```
            # logging
```

```
            losses.append(loss)
```

```
    return weights, losses
```

```
    # compute gradient
```

```
    def compute_gradient_w(X_features, y, y_hat):
```

```
        dl_dweights = [2*xi*(y_hat-y) for xi in X_features]
```

```
        return dl_dweights
```


Exercise 4

```
▶ def implement_linear_regression(X_feature, y_output, epoch_max = 50, lr = 1e-5):
```

```
    losses = []
```

```
    weights = initialize_params()
```

```
    N = len(y_output)
```

```
    for epoch in range(epoch_max):
```

```
        print("epoch", epoch)
```

```
        for i in range(N):
```

```
            # get a sample - row i
```

```
            features_i = X_feature[i]
```

```
            y = sales_data[i]
```

```
            # compute output
```

```
            y_hat = predict(features_i, weights)
```

```
            # compute loss
```

```
            loss = compute_loss(y, y_hat)
```

```
            # compute gradient w1, w2, w3, b
```

```
            dl_dweights = compute_gradient_w(features_i, y, y_hat)
```

```
            # update parameters
```

```
            weights = update_weight(weights, dl_dweights, lr)
```

```
            # logging
```

```
            losses.append(loss)
```

```
    return weights, losses
```

```
    # update weights
```


```
    def update_weight(weights, dl_dweights, lr):
```

```
        weights = [w - lr*dw for w, dw in zip(weights, dl_dweights)]
```

```
        return weights
```

Exercise 4

```
X,y = prepare_data('advertising.csv')  
W,L = implement_linear_regression(X,y)  
plt.plot(L[-100:])  
plt.xlabel("#iteration")  
plt.ylabel("MSE loss")  
plt.show()
```



Outline

➤ **Linear Regression Review**

➤ **Exercise 1**

➤ **Exercise 2**


➤ **Exercise 3**

➤ **Exercise 4**

➤ **Exercise 5**

➤ **Other Discussions**





Exercise 5

Bài tập 5 (Tìm hiểu kỹ thuật Feature Scaling thông qua min and max)): Ở bài tập này các bạn cần

12

AI VIETNAM

aivietnam.edu.vn

phải chuẩn hoá dữ liệu đầu vào X trước thuật nhằm tăng tốc độ hội tụ của giải thuật linear regression. Yêu cầu của bài tập là các bạn cần chuẩn hoá data theo MinMax scale $X_{new} = \frac{X_{old} - X_{min}}{X_{max} - X_{min}}$ trước khi đưa vào huấn luyện data. Nhiệm vụ của bạn ở bài tập này là cần implement hàm **min_max_scaling()** để chuẩn hoá dữ liệu input X như bên dưới:

Exercise 5

```
def prepare_data(file_name_dataset):
    data = np.genfromtxt(file_name_dataset, delimiter=',', skip_header=1).tolist()

    # get tv (index=0)
    tv_data = get_column(data, 0)

    # get radio (index=1)
    radio_data = get_column(data, 1)

    # get newspaper (index=2)
    newspaper_data = get_column(data, 2)

    # get sales (index=3)
    sales_data = get_column(data, 3)


    # scale data (only for features)
    # remember to scale input features in inference, therefore, we need to save max, min and mean values
    (tv_data, radio_data, newspaper_data), (max_data_1, max_data_2, max_data_3, min_data_1, min_data_2, min_data_3)
    | = min_max_scaling(tv_data, radio_data, newspaper_data)

    # building X input and y output for training
    # Create list of features for input
    X = [[1, x1, x2, x3] for x1, x2, x3 in zip(tv_data, radio_data, newspaper_data)]
    y = sales_data
    return X, y
```


Exercise 5

```
def min_max_scaling(data1, data2, data3):  
    max_data_1 = max(data1)  
    max_data_2 = max(data2)  
    max_data_3 = max(data3)  
  
    min_data_1 = min(data1)  
    min_data_2 = min(data2)  
    min_data_3 = min(data3)  
  
    data1 = [(x - min_data_1) / (max_data_1 - min_data_1) for x in data1]  
    data2 = [(x - min_data_2) / (max_data_2 - min_data_2) for x in data2]  
    data3 = [(x - min_data_3) / (max_data_3 - min_data_3) for x in data3]  
  
    return (data1, data2, data3), (max_data_1, max_data_2, max_data_3, min_data_1, min_data_2, min_data_3)
```


Exercise 5



Before Feature Scaling



After Feature Scaling

Outline

➤ **Linear Regression Review**

➤ **Exercise 1**

➤ **Exercise 2**

➤ **Exercise 3**

➤ **Exercise 4**

➤ **Exercise 5**


➤ **Other Discussions**



How to measure the accuracy of the system

X

Y



TV	Radio	Newspaper	Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9
8.7	48.9	75	7.2
57.5	32.8	23.5	11.8
120.2	19.6	11.6	13.2
8.6	2.1	1	4.8
199.8	2.6	21.2	15.6
66.1	5.8	24.2	12.6
214.7	24	4	17.4
23.8	35.1	65.9	9.2
97.5	7.6	7.2	13.7
204.1	32.9	46	19
195.4	47.7	52.9	22.4
67.8	36.6	114	12.5
281.4	39.6	55.8	24.4
69.2	20.5	18.3	11.3
147.3	23.9	19.1	14.6
218.4	27.7	53.4	18
237.4	5.1	23.5	17.5

Training

Strategy:


➤ Train: 200 samples

How to measure the accuracy of the system

X

Y

TV	Radio	Newspaper	Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.2	58.5	16.5
180.8	42.8	57.5	17.9
8.7	32.8	1	7.2
57.5	19.6	21.2	11.8
120.2	2.1	24.2	13.2
8.6	2.6	4	4.8
199.8	5.8	65.9	15.6
66.1	24	7.2	12.6
214.7	35.1	46	17.4
23.8	7.6	52.9	9.2
97.5	32.9	114	13.7
204.1	47.7	55.8	19
195.4	36.6	18.3	22.4
67.8	39.6	19.1	12.5
281.4	20.5	53.4	24.4
69.2	23.9	23.5	11.3
147.3	27.7		14.6
218.4	5.1		18
237.4			17.5



Testing

Strategy: $n = 200$

- Train: 200 samples
- Test: 200 samples

100% dataset for
training
100% dataset for


$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

How to measure the accuracy of the system

X

Y

TV	Radio	Newspaper	Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.2	58.5	16.5
180.8	41.2	58.5	17.9
8.7	41.2	58.5	7.2
57.5	32.8	58.5	11.8
120.2	19.6	58.5	13.2
8.6	2.1	58.5	4.8
199.8	2.6	58.5	15.6
66.1	5.8	58.5	12.6
214.7	24	58.5	17.4
23.8	35.1	58.5	9.2
97.5	7.6	58.5	13.7
204.1	32.9	58.5	19
195.4	47.7	58.5	22.4
67.8	36.6	58.5	12.5
281.4	39.6	58.5	24.4
69.2	20.5	58.5	11.3
147.3	23.9	58.5	14.6
218.4	27.7	58.5	18
237.4	5.1	58.5	17.5



Testing

Strategy:

- Train: 160 samples
- Test: 40 samples

80% dataset for training
20% dataset for testing
Train set \neq Test set



$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

How to measure the accuracy of the system

X

Y

TV	Radio	Newspaper	Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.2	58.5	16.5
180.8	42.8	53.5	17.9
8.7	43.8	54.5	7.2
57.5	32.8	55.5	11.8
120.2	19.6	56.5	13.2
8.6	2.1	57.5	4.8
199.8	2.6	58.5	15.6
66.1	5.8	59.5	12.6
214.7	24	60.5	17.4
23.8	35.1	61.5	9.2
97.5	7.6	62.5	13.7
204.1	32.9	63.5	19
195.4	47.7	64.5	22.4
67.8	36.6	65.5	12.5
281.4	39.6	66.5	24.4
69.2	20.5	67.5	11.3
147.3	23.9	68.5	14.6
218.4	27.7	69.5	18
237.4	5.1	70.5	17.5



Testing

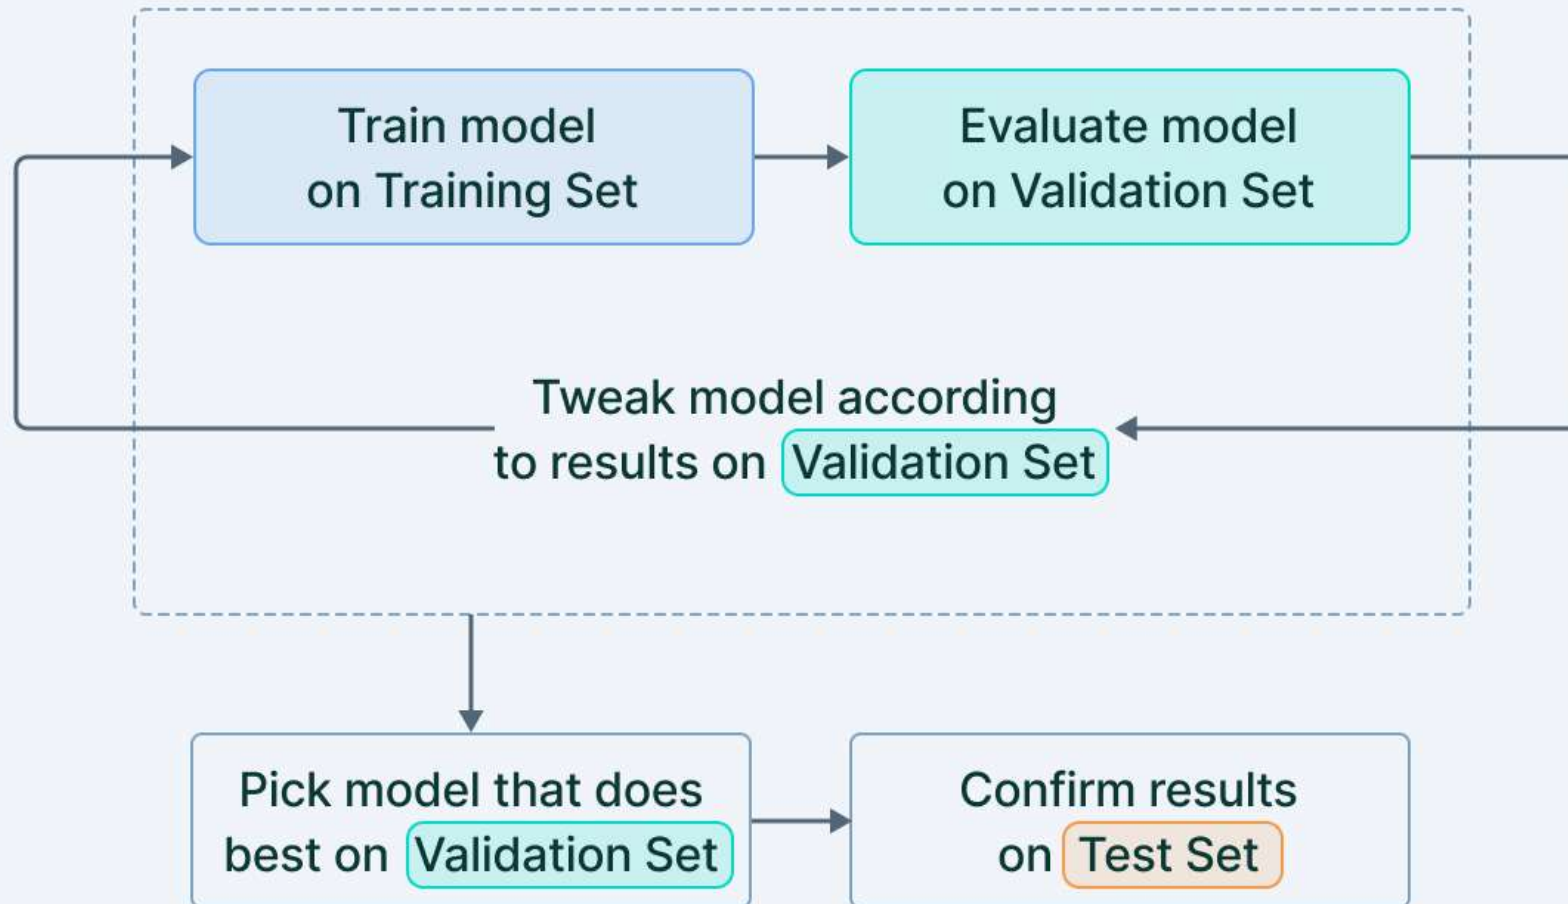
Strategy:

- Train: 160 samples
- Test: 40 samples

60% dataset for
training
20% dataset for

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

Train/Validation/Test



Exercise 5






Training data



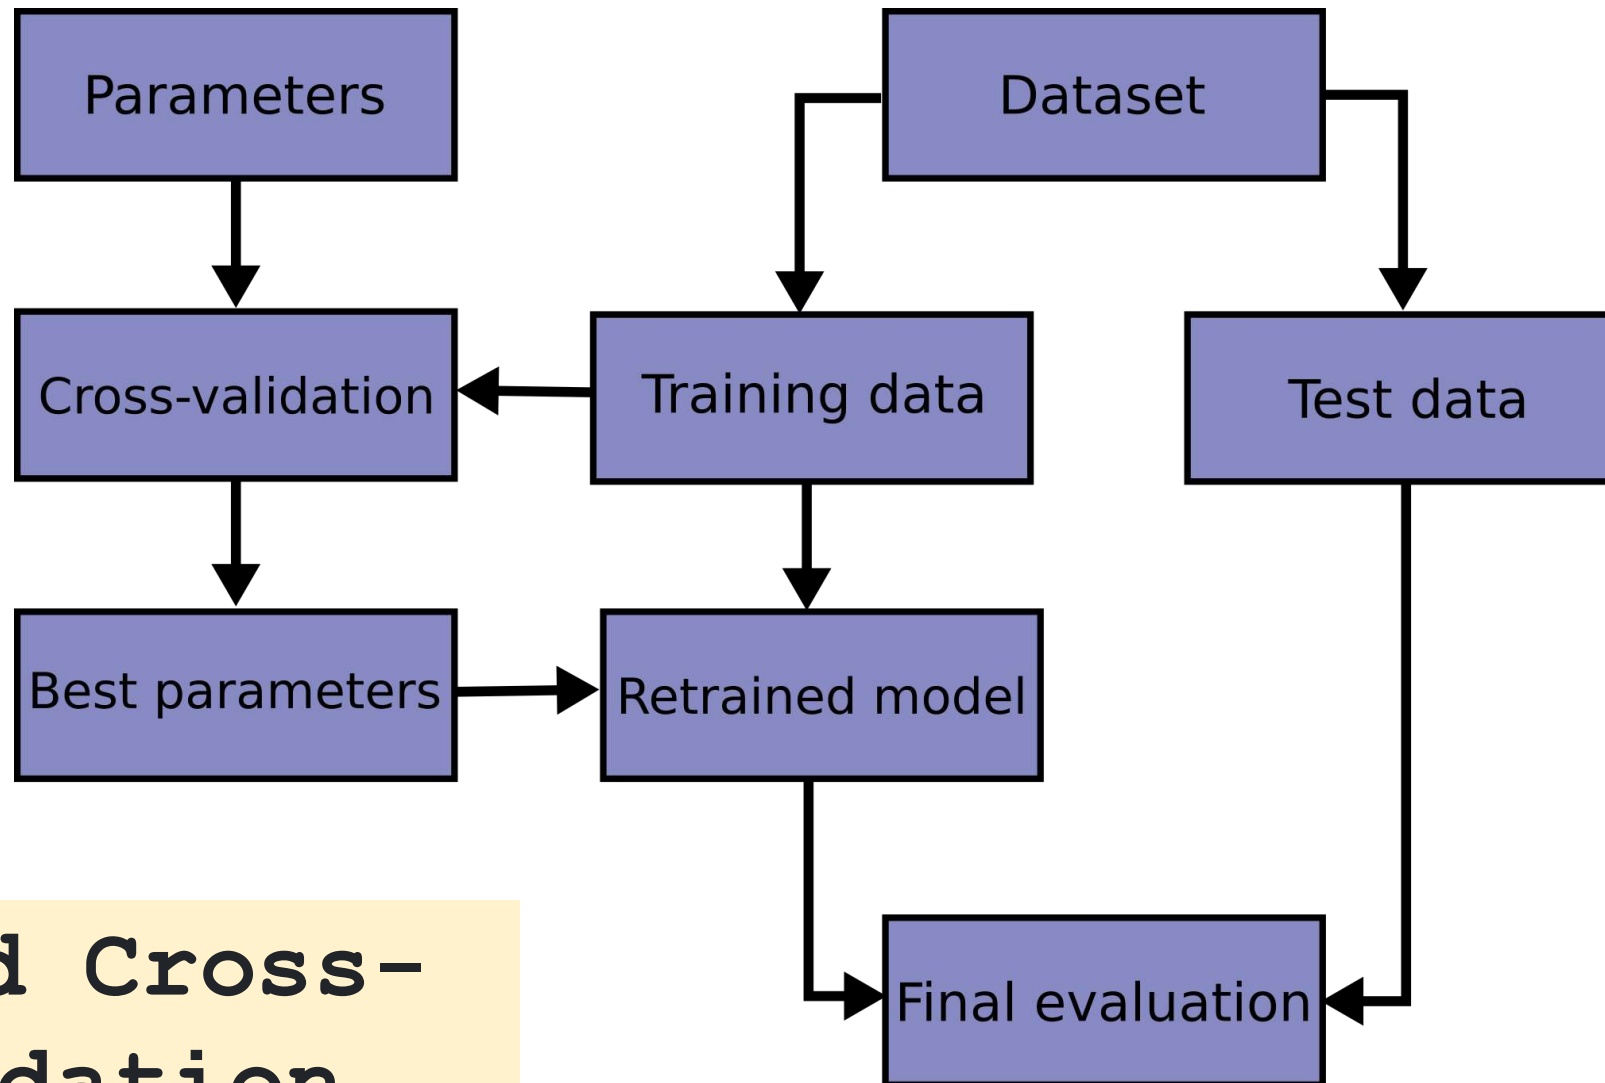
Validation data



Test data



Exercise 5



K Fold Cross-validation

Exercise 5

